

Project Documentation: Hadoop + Hive Pipeline for MIMIC-III Data

Project Overview

This project demonstrates how I set up a simple Hadoop and Hive big data pipeline using Docker. The goal is to store and query MIMIC-III clinical data (in Parquet format) using Hive on top of HDFS.

Step-by-Step Implementation

1. Clone the Docker Hadoop-Spark Project

I used an existing GitHub repository to quickly spin up a Hadoop + Hive environment:

```
yasmin@Yasmin MINGW64 /D/GitHub (master)
$ git clone https://github.com/Marcel-Jan/docker-hadoop-spark.git
cd docker-hadoop-spark
```

2. Add Persistent Storage for Hive Metastore

To make sure Hive external tables don't get lost when restarting the containers, I modified the docker-compose.yml file to persist the PostgreSQL database used by Hive Metastore:

```
156     hive-metastore-postgresql:
157       image: bde2020/hive-metastore-postgresql:2.3.0
158       container_name: hive-metastore-postgresql
159       ports:
160         - "5432:5432"
161       volumes:
162         - hive_pgdata:/var/lib/postgresql/data
163     presto-coordinator:
```

3. Copy Data Files into Namenode

I copied the cleaned .parquet files from my local machine into the namenode container using docker cp:

```
D:\GitHub\docker-hadoop-spark>docker cp "D:\iti\big_data\data_cleaned\diagnosis_full.p
arquet" namenode:/diagnosis_full.parquet
Successfully copied 60.4kB to namenode:/diagnosis_full.parquet

D:\GitHub\docker-hadoop-spark>docker cp "D:\iti\big_data\data_cleaned\icustays.parquet
" namenode:/icustays.parquet
Successfully copied 17.4kB to namenode:/icustays.parquet

D:\GitHub\docker-hadoop-spark>docker cp "D:\iti\big_data\data_cleaned\patients.parquet
" namenode:/patients.parquet
Successfully copied 12.3kB to namenode:/patients.parquet
```

(I did this for the diagnosis_full.parquet file as well)

4. Upload Data to HDFS

First I created a directory for every file:

```
root@240339a68101:/# hdfs dfs -mkdir -p /user/hive/warehouse/admissions
root@240339a68101:/# hdfs dfs -mkdir -p /user/hive/warehouse/diagnosis_full
root@240339a68101:/# hdfs dfs -mkdir -p /user/hive/warehouse/icustays
root@240339a68101:/# hdfs dfs -mkdir -p /user/hive/warehouse/patients
root@240339a68101:/# |
```

Then I uploaded the files:

```
root@240339a68101:/# hdfs dfs -put admissions.parquet /user/hive/warehouse/admissions/
2025-05-18 15:03:38,623 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteH
ostTrusted = false
root@240339a68101:/# hdfs dfs -put diagnosis_full /user/hive/warehouse/diagnosis_full/
put: 'diagnosis_full': No such file or directory
root@240339a68101:/# hdfs dfs -put diagnosis_full.parquet /user/hive/warehouse/diagnosis_full/
2025-05-18 15:04:25,518 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteH
ostTrusted = false
root@240339a68101:/# hdfs dfs -put icustays.parquet /user/hive/warehouse/icustays/
2025-05-18 15:04:53,987 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteH
ostTrusted = false
root@240339a68101:/# hdfs dfs -put patients.parquet /user/hive/warehouse/patients/
2025-05-18 15:05:12,128 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteH
ostTrusted = false
root@240339a68101:/# |
```

5. Create External Tables in Hive

I went inside hive-server container with the following code: “docker exec -it hive-server bash”, then I got connected to hive by just typing “hive”.

I created a database named mimic using the following code:

Create database mimic;

Use mimic;

Now we can start creating our tables.

Admissions Table:

```
0: jdbc:hive2://localhost:10000> create external table admissions(
. . . . .> row_id INT,
. . . . .> subject_id INT,
. . . . .> admittime TIMESTAMP,
. . . . .> disctime TIMESTAMP,
. . . . .> deathtime TIMESTAMP,
. . . . .> admission_type STRING,
. . . . .> admission_location STRING,
. . . . .> discharge_location STRING,
. . . . .> insurance STRING,
. . . . .> language STRING,
. . . . .> religion STRING,
. . . . .> ethnicity STRING,
. . . . .> edregtime TIMESTAMP,
. . . . .> edouttime TIMESTAMP,
. . . . .> diagnosis STRING,
. . . . .> hospital_expire_flag INT,
. . . . .> has_chartevents_data INT)
. . . . .> stored as parquet
. . . . .> location '/user/hive/warehouse/admissions/';
No rows affected (1.016 seconds)
```

Diagnosis_full Table:

```
0: jdbc:hive2://localhost:10000> create external table diagnosis_full(
. . . . .> row_id INT,
. . . . .> subject_id INT,
. . . . .> hadm_id INT,
. . . . .> seq_num INT,
. . . . .> icd9_code STRING,
. . . . .> short_title STRING,
. . . . .> long_title STRING)
. . . . .> stored as parquet
. . . . .> location '/user/hive/warehouse/admissions/';
No rows affected (0.043 seconds)
0: jdbc:hive2://localhost:10000> |
```

icustays Table:

```
0: jdbc:hive2://localhost:10000> create external table icustays(
. . . . .> row_id INT,
. . . . .> subject_id INT,
. . . . .> hadm_id INT,
. . . . .> icustay_id INT,
. . . . .> dbsource STRING,
. . . . .> first_careunit STRING,
. . . . .> last_careunit STRING,
. . . . .> first_wardid INT,
. . . . .> last_wardid INT,
. . . . .> intime TIMESTAMP,
. . . . .> outtime TIMESTAMP)
. . . . .> stored as parquet
. . . . .> location '/user/hive/warehouse/icustays/';
No rows affected (0.046 seconds)
```

Patients Table:

```
0: jdbc:hive2://localhost:10000> create external table patients(
. . . . .> row_id INT,
. . . . .> subject_id INT,
. . . . .> gender STRING,
. . . . .> dob TIMESTAMP,
. . . . .> dod TIMESTAMP,
. . . . .> dod_hosp TIMESTAMP,
. . . . .> dod_ssn TIMESTAMP,
. . . . .> expire_flag INT)
. . . . .> stored as parquet
. . . . .> location '/user/hive/warehouse/patients/';
No rows affected (0.036 seconds)
```

Now that I have my tables ready, I can start querying.

6. Querying

1st Question: Average Length per Diagnosis:

I ran the following query:

```
SELECT
```

```
AVG(UNIX_TIMESTAMP(disctime) - UNIX_TIMESTAMP(admittime)) / 86400 AS avg_los,
```

```
diagnosis
```

FROM admissions

GROUP BY diagnosis;

Output:

-c0	diagnosis
6.960267452697917E9	MITRAL REGURGITATION;CORONARY ARTERY DISEASE\CORONARY ARTERY BYPASS GRAFT WITH MVR ? MITRAL VALVE REPLACEMENT /SDA
5.843995963892361E9	ABDOMINAL PAIN
6.031651308315278E9	ABSCESS
6.291175230030556E9	ACUTE CHOLANGITIS
4.65877899223125E9	ACUTE CHOLECYSTITIS
4.2521854102701387E9	ACUTE PULMONARY EMBOLISM
4.859564050406944E9	ACUTE RESPIRATORY DISTRESS SYNDROME;ACUTE RENAL FAILURE
6.434881724086111E9	ACUTE SUBDURAL HEMATOMA
6.102058974136806E9	ALCOHOLIC HEPATITIS
6.575430995146528E9	ALTERED MENTAL STATUS
5.522424889322222E9	AROMEGLEY;BURKITT'S LYMPHOMA
6.006846979409722E9	ASTHMA/COPD FLARE
5.5419984642878475E9	ASTHMA;CHRONIC OBST PULM DISEASE
5.131158313216666E9	BASAL GANGLIN BLEED
6.644302698711805E9	BRADYCARDIA
7.087537498011806E9	BRAIN METASTASES
5.508939508186111E9	CELLULITIS
5.456568148111111E9	CEREBROVASCULAR ACCIDENT
4.646582029461805E9	CHEST PAIN
4.32673054651875E9	CHEST PAIN/ CATH
6.342269802865278E9	CHOLANGITIS
5.992340834885555E9	CHOLECYSTITIS
4.2740805432375E9	CHRONIC MYELOGENOUS LEUKEMIA;TRANSFUSION REACTION
5.991224376060417E9	CONGESTIVE HEART FAILURE
6.49319850479167E9	CORONARY ARTERY DISEASE\CORONARY ARTERY BYPASS GRAFT /SDA
5.398393418642361E9	CRITICAL AORTIC STENOSIS/HYPOTENSION
5.5528412472409725E9	ELEVATED LIVER FUNCTIONS;S/P LIVER TRANSPLANT
5.853125898182292E9	ESOPHAGEAL CA/SDA
6.933656560697917E9	ESOPHAGEAL CANCER/SDA
5.518104255373611E9	FACIAL NUMBNESS
5.087746922915972E9	FAILURE TO THRIVE
5.743940510158333E9	FEVER
5.868877620319445E9	FEVER;URINARY TRACT INFECTION
5.440164340893056E9	GASTROINTESTINAL BLEED
6.049978381068055E9	HEADACHE
4.980343778125E9	HEPATIC ENCEP
4.943559936061111E9	HEPATITIS B
5.661966595278472E9	HUMERAL FRACTURE
6.184578424486806E9	HYPOGLYCEMIA
4.754498097264584E9	HYPONATREMIA;URINARY TRACT INFECTION
6.024048787569792E9	HYPOTENSION
6.60507725745E9	HYPOTENSION, RENAL FAILURE
6.6329145392125E9	HYPOTENSION;TELEMETRY
7.227850126282639E9	HYPOTENSION;UNRESPONSIVE
6.189969969252084E9	INFERIOR MYOCARDIAL INFARCTION\CATH
7.013836090354167E9	LEFT HIP FRACTURE
7.062474942697917E9	LEFT HIP OA/SDA
5.474914301064583E9	LIVER FAILURE
5.70101762226875E9	LOWER GI BLEED
4.330879482113889E9	LUNG CANCER;SHORTNESS OF BREATH
6.119335183226389E9	MEDIASTINAL ADENOPATHY
5.767880970316667E9	METASTATIC MELANOMA;BRAIN METASTASIS
4.450640443247916E9	METASTIC MELANOMA;ANEMIA
6.517908164441667E9	MI CHF
4.850576683166667E9	NON SMALL CELL CANCER;HYPOXIA
4.860876141398611E9	OVERDOSE
4.49470938734375E9	PERICARDIAL EFFUSION
4.312660985345139E9	PLEURAL EFFUSION
6.305283676407553E9	PNEUMONIA
6.060507357419444E9	PNEUMONIA/HYPOGLCEMIA/SYNCOPE
5.767732823204861E9	PNEUMONIA;TELEMETRY

Question 2: Distribution of ICU readmissions:

Query:

SELECT COUNT(subject_id), subject_id

FROM icustays

GROUP BY subject_id

HAVING COUNT(subject_id) > 1;

Output:

```
0: jdbc:hive2://localhost:10000> select count(subject_id), subject_id from icustays group by subject_id having count(subject_id)>1
. . . . .>;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
+-----+-----+
|_c0|subject_id|
+-----+-----+
|2|10059|
|3|10088|
|2|10094|
|2|10117|
|2|10119|
|4|10124|
|2|40124|
|2|40177|
|2|40304|
|2|40310|
|2|41795|
|15|41976|
|2|42135|
|2|42281|
|2|42346|
|2|43735|
|2|43746|
|2|43881|
|3|44083|
+-----+-----+
10 rows selected (1.3 seconds)
0: jdbc:hive2://localhost:10000> |
```

3rd Question: Mortality rates by demographic groups:

Query:

SELECT

(SUM(CASE WHEN hospital_expire_flag = 1 THEN 1 ELSE 0 END) * 100 / COUNT(*)) AS mortality_rate,

ethnicity

FROM admissions

GROUP BY ethnicity;

```
0: jdbc:hive2://localhost:10000> select (sum(case when hospital_expire_flag = 1 then 1 else 0 end)*100 / count(*)) as mortality_rate, ethnicity
. . . . .> from admissions
. . . . .> group by ethnicity;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
+-----+-----+
|mortality_rate|ethnicity|
+-----+-----+
|58.0|AMERICAN INDIAN/ALASKA NATIVE FEDERALLY RECOGNIZED TRIBE|
|58.0|ASIAN|
|28.571428571428573|BLACK/AFRICAN AMERICAN|
|5.882352941176471|HISPANIC OR LATINO|
|66.66666666666667|OTHER|
|58.0|UNKNOWN/NOT SPECIFIED|
|31.3953488372893|WHITE|
+-----+-----+
7 rows selected (1.312 seconds)
0: jdbc:hive2://localhost:10000> |
```

YARN Web UI Setup

To access the YARN ResourceManager on the web, I edited the docker-compose.yml file and added the ports part:

```

76   resourcemanager:
77     image: bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8
78     container_name: resourcemanager
79     restart: always
80     environment:
81       SERVICE_PRECONDITION: "namenode:9000 namenode:9870 datanode:9864"
82     ports:
83       - "8088:8088"

```

I also checked the yarn-site.xml config inside the resourcemanager container to make sure it's binding to 0.0.0.0:8088.

Hue Web UI Setup for Hive

To install Hue and connect it with Hive via Web UI, I added the following to my docker-compose.yml:

```

hue:
  image: gethue/hue:latest
  container_name: hue
  hostname: hue
  ports:
    - "8888:8888"
  depends_on:
    - hive-server
    - hive-metastore
    - hive-metastore-postgresql
    - database
  volumes:
    - .\data\hue\hue-overrides.ini:/usr/share/hue/desktop/conf/z-hue-overrides.ini
  environment:
    - DATABASE_ENGINE=mysql
    - DATABASE_NAME=hue
    - DATABASE_USER=root
    - DATABASE_PASSWORD=secret
    - DATABASE_HOST=database
    - DATABASE_PORT=3306
    - HIVE_SERVER_HOST=hive-server
    - HIVE_SERVER_PORT=10000

```

And for the database:

```

database:
  image: mysql:5.7
  container_name: database
  hostname: database
  ports:
    - "33061:3306"
  environment:
    MYSQL_ROOT_PASSWORD: secret
    MYSQL_DATABASE: hue
    MYSQL_USER: root
    MYSQL_PASSWORD: secret
  volumes:
    - ./data/mysql/data:/var/lib/mysql
    - ./data/init.sql:/docker-entrypoint-initdb.d/init.sql
  command: >
    mysqld --innodb-flush-method=O_DSYNC
           --innodb-use-native-aio=OFF
           --init-file /docker-entrypoint-initdb.d/init.sql

```

Currently, I can see the tables in the Hue Web UI, but queries result in a "database is locked" error. Likely because Hue is not properly connected to the MySQL backend.

MapReduce Job: Calculate Average Age

1. Create a Table for MapReduce

I created a table formatted in a way that map reduce understands

```

hive> CREATE EXTERNAL TABLE mimic.patients_for_mr (
>   subject_id INT,
>   dob STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE
> LOCATION '/user/hive/warehouse/patients_for_mr/patients_for_mr';
OK

```

2. Insert the Data


```

Time taken: 3.611 seconds
hive> INSERT OVERWRITE TABLE mimic.patients_for_mr
> SELECT subject_id, CAST(dob AS STRING) FROM mimic.patients;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = root_20250522214249_6c159286-fd75-44ec-b398-c622886241a8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Job running in-process (local Hadoop)
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
2025-05-22 21:42:52 748 Stage-1 map = 100%, reduce = 0%
Ended Job = job_local2060275424_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://namenode:9000/user/hive/warehouse/patients_for_mr/patients_for_mr/.hive-staging_hive_2025-05-22_21-42-49_703_5070404082218854174-1/-ext-10
000
Loading data to table mimic.patients_for_mr
MapReduce Jobs Launched:
Stage-Stage-1:  HDFS Read: 10231 HDFS Write: 2680 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 3.611 seconds
hive>

```

2. Java Setup for MapReduce

I made sure first that java is available in the namenode container then I created 2 directories; 1 for the mapper and the other for the reducer.

Inside the namenode:

```

D:\GitHub\docker-hadoop-spark>docker exec -it namenode bash
root@d004fad20896:/# hadoop version
Hadoop 3.2.1
Source code repository https://gitbox.apache.org/repos/asf/hadoop.git -r b3cbbb467e
22ea829b3808f4b7b01d07e0bf3842
Compiled by rohithsharmaks on 2019-09-10T15:56Z
Compiled with protoc 2.5.0
From source with checksum 776eaf9eee9c0ffc370bcbcb1888737
This command was run using /opt/hadoop-3.2.1/share/hadoop/common/hadoop-common-3.2.
1.jar
root@d004fad20896:/# javac -version
javac 1.8.0_232
root@d004fad20896:/# mkdir -p ~/avg-age-job/src ~/avg-age-job/build
root@d004fad20896:/# |

```

4. Compile and Package the Job

After putting the code into the files I created:

```

root@d004fad20896:/# javac -classpath `hadoop classpath` -d ~/avg-age-job/build ~/a
vg-age-job/src/AverageAge.java

```

5. Run the Job

```

root@d004fad20896:/# javac -classpath `hadoop classpath` -d ~/avg-age-job/build ~/a
vg-age-job/src/AverageAge.java
root@d004fad20896:/# jar -cvf ~/avg-age-job/average-age.jar -C ~/avg-age-job/build/
.
added manifest
adding: AverageAge.class(in = 1450) (out= 797)(deflated 45%)
adding: AverageAge$AgeMapper.class(in = 2305) (out= 1021)(deflated 55%)
adding: AverageAge$AvgReducer.class(in = 1782) (out= 771)(deflated 56%)
root@d004fad20896:/# |

```

Summary

I now have a working Dockerized Hadoop + Hive environment:

- **Data** stored in HDFS (as Parquet)
 - **Hive** external tables mapped to that data
 - **Queries** run directly in Hive
 - **Hue** as a GUI (with issues to fix)
 - **YARN** accessible via browser
 - **MapReduce job** successfully compiled and executed
-

Tools Used

- Docker
 - Hadoop (HDFS)
 - Hive + Hive Metastore
 - PostgreSQL / MySQL
 - Parquet files
 - Hue Web UI
 - Java (MapReduce)
-

Credits

Based on: <https://github.com/Marcel-Jan/docker-hadoop-spark>

Problems Faced

1. Parquet-Hive Compatibility Issue

When I first tried to load the Parquet files into Hive, I encountered an error related to data types — Hive didn't support the timestamp format in the Parquet files generated by Pandas.

I initially tried to cast everything to string and later convert it to TIMESTAMP, but this didn't work properly. Hive kept showing me that the field was not compatible with timestamp. I wanted to keep datetime columns in proper format, so I fixed the problem by modifying my Parquet export code in Python:

I used the pyarrow engine and added:

```
use_deprecated_int96_timestamps=True
```

This allowed Hive to recognize the timestamp columns correctly.
